

Chapter 4: Classification & Prediction

▶ 4.1 Basic Concepts of Classification and Prediction

4.1.1 Definition

4.1.2 Classification vs. Prediction

4.1.3 Classification Steps

4.1.4 Issues of Classification and Prediction

▶ 4.2 Decision Tree Induction

4.2.1 The Algorithm

4.2.2 Attribute Selection Measures

4.2.3 Tree Pruning

4.2.4 Scalability and Decision Tree Induction

▶ 4.3 Bayes Classification Methods

▶ 4.4 Rule Based Classification

▶ 4.5 Lazy Learners

▶ 4.6 Prediction

▶ 4.7 How to Evaluate and Improve Classification

4.1.1 Definition

- ▶ Classification is also called **Supervised Learning**

- ▶ **Supervision**


- The training data (observations, measurements, etc) are used to learn a classifier

- The training data are **labeled** data

- New data (**unlabeled**) are classified

Using the training data

Training data		
Age	Income	Class label
27	28K	Budget-Spenders
35	36K	Big-Spenders
65	45K	Budget-Spenders



- ▶ **Principle**


- Construct models (functions) based on some training examples
 - Describe and distinguish classes or concepts for future prediction
 - Predict some unknown class labels

4.1.2 Classification vs. Prediction

► Classification

- Predicts categorical class labels (discrete or nominal)
- Use labels of the training data to classify new data

► Example




- A model or classifier is constructed to predict **categorical labels** such as “safe” or “risky” for a loan application data.

► Prediction

- Models continuous-valued functions, i.e., predicts unknown or missing values

► Example

- A marketing manager would like to predict how much a given customer will spend during a sale



- Unlike classification, it provides ordered values
- **Regression** analysis is used for prediction
- Prediction is a short name for **numeric prediction**

4.1.3 Classification Steps (1/2)

There are two main steps in classification

► **Step1: Model Construction (learning step, or training step)**

→ Construct a classification model based on **training data**

→ **Training data**

- A set of tuples
- Each tuple is assumed to belong to a predefined class
- Labeled data (ground truth)

→ **How a classification model looks like?**

A classification model can be represented by one of the following forms:

- Classification rules
- Decision trees
- Mathematical formulae

4.1.3 Classification Steps (2/2)


► Step2: Model Usage

Before using the model, we first need to test its accuracy

→ Measuring model accuracy

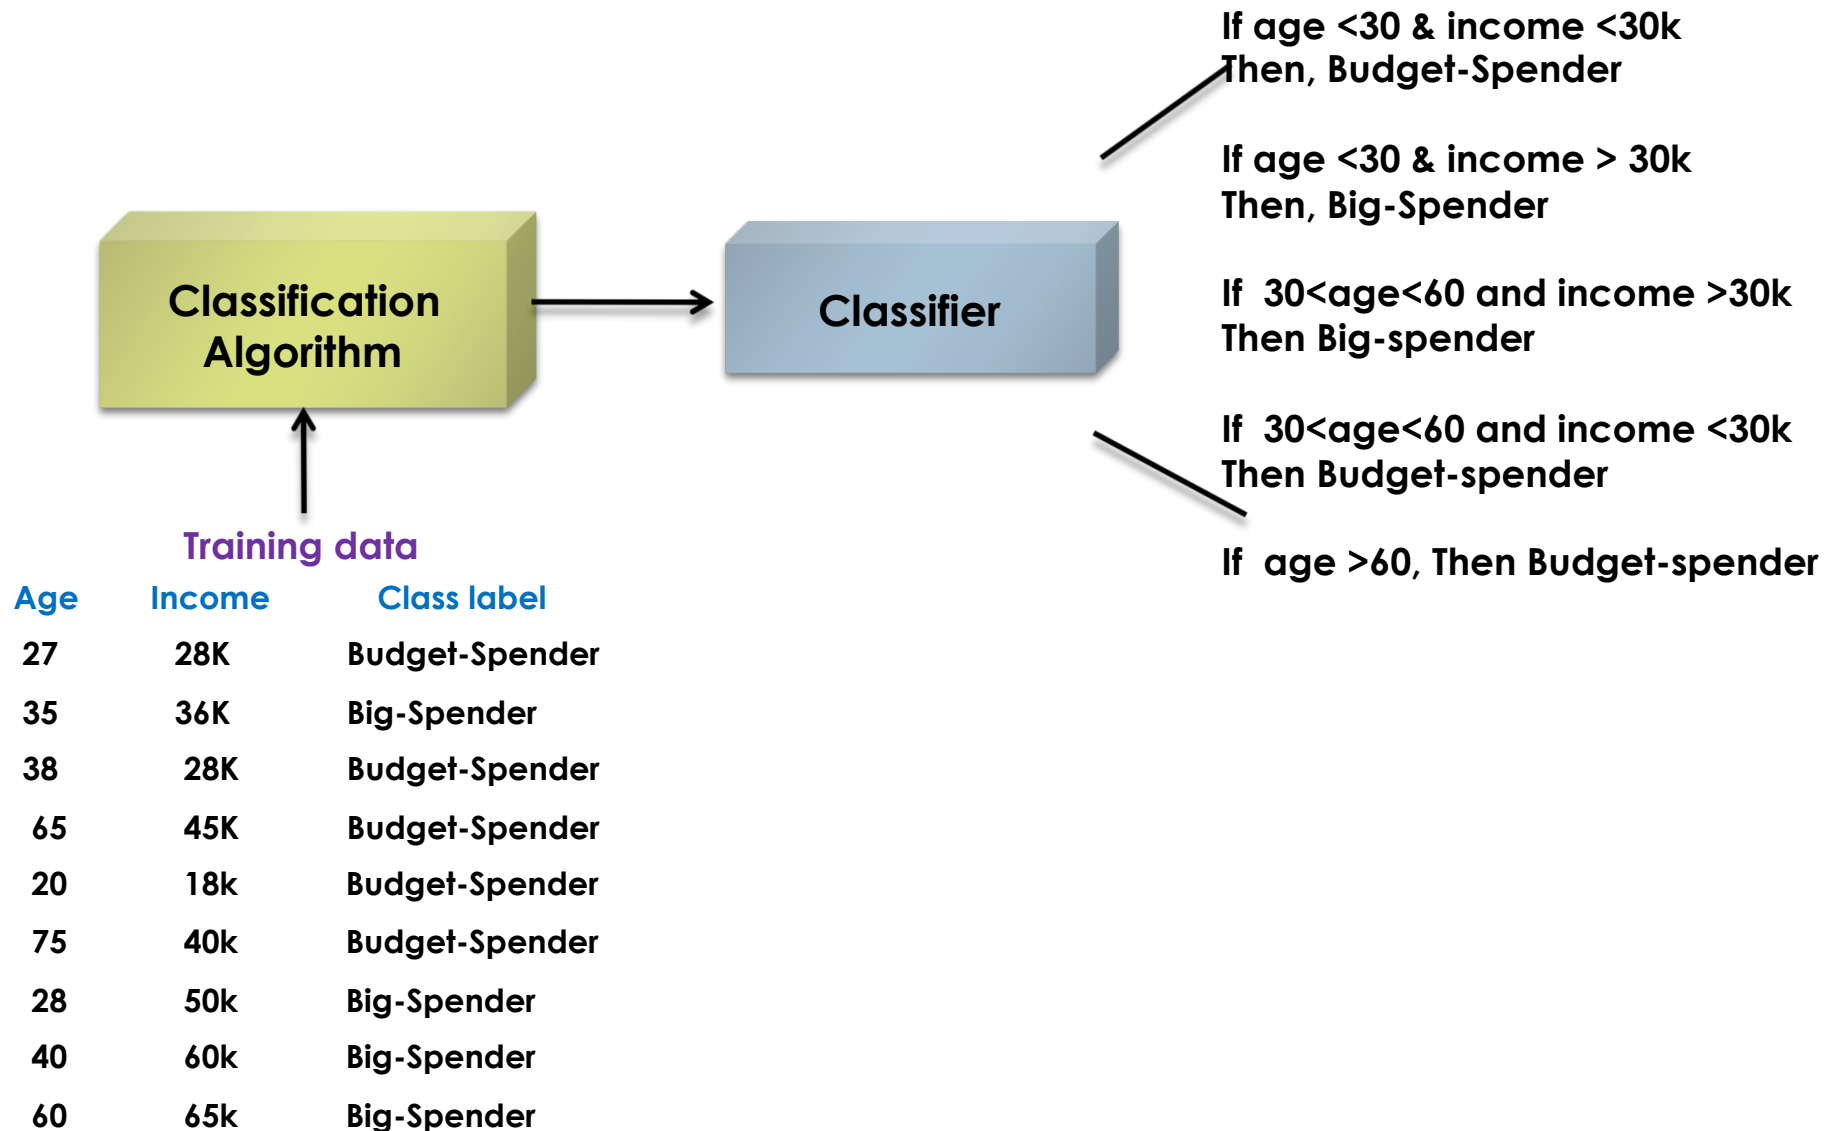
- To measure the accuracy of a model we need **test data**
- Test data is similar in its structure to training data (labeled data)
- **How to test?**

The known label of test sample is compared with the classified result from the model




- **Accuracy rate** is the percentage of test set samples that are correctly classified by the model
 - **Important:** test data should be independent of training set, otherwise over-fitting will occur
- **Using the model:** If the accuracy is acceptable, use the model to **classify data** tuples whose class labels are not known

Model Construction



Model Usage


1-Test the classifier



2-If acceptable accuracy

Unlabeled data

Age	Income
18	28K
37	40K
60	45K
40	36k



Classified data

Age	Income	Class label
18	28K	Budget-Spenders
37	40K	Big-Spenders
60	45K	Budget-Spenders
40	36k	Budget-Spenders

4.1.4 Issues of Classification & Prediction

Data Preparation

▶ Data cleaning

- Perform a preprocessing step to reduce noise and handle missing values
- How to handle missing values?
 - E.g., replacing a missing value with the most commonly occurring value for that attribute, the most probable value based on statistics (prediction) (example)

▶ Relevance analysis (feature selection)

- Remove irrelevant or redundant attributes

▶ Data transformation and reduction

- Generalize data to (higher concepts, discretization)
- Normalizing attribute values (income vs. binary attributes)
- Reduce the dimensionality of the data

4.1.4 Issues of Classification & Prediction

Evaluating Classification Methods

▶ Accuracy

- classifier accuracy: the ability of a classifier to predict class labels
- predictor accuracy: how close is the predicted value from true one.

▶ Speed

- time to construct the model (training time)
- time to use the model (classification/prediction time)

▶ Robustness

- handling noise and missing values

▶ Scalability

- efficiency in disk-resident databases

▶ Interpretability

- Level of understanding and insight provided by the model

Summary of section 4.1.1

- ▶ **Classification** predicts class labels
- ▶ Numeric **prediction** models continued-valued functions
- ▶ Two steps of classification: **1) Training**
2) Testing and using
- ▶ Data cleaning and Evaluation are the main issues of classification and prediction

Chapter 4: Classification & Prediction

▶ 4.1 Basic Concepts of Classification and Prediction

4.1.1 Definition

4.1.2 Classification vs. Prediction

4.1.3 Classification Steps

4.1.4 Issues of Classification and Prediction

▶ 4.2 Decision Tree Induction

4.2.1 The Algorithm

4.2.2 Attribute Selection Measures

4.2.3 Tree Pruning

4.2.4 Scalability and Decision Tree Induction

▶ 4.3 Bayes Classification Methods

▶ 4.4 Rule Based Classification


▶ 4.5 Lazy Learners

▶ 4.6 Prediction

▶ 4.7 How to Evaluate and Improve Classification

4.2 Decision Tree Induction

- ▶ **Decision tree induction** is the learning of decision trees from class-labeled training tuples
- ▶ A decision tree is a flowchart-like tree structure
 - **Internal nodes** (non leaf node) denotes a test on an attribute
 - **Branches** represent outcomes of tests
 - **Leaf nodes** (terminal nodes) hold class labels
 - **Root node** is the topmost node



A decision tree indicating whether a customer is likely to purchase a computer

Class-label Yes: The customer is likely to buy a computer

Class-label no: The customer is unlikely to buy a computer

4.2 Decision Tree Induction


► How are decision trees used for classification?

- The attributes of a tuple are tested against the decision tree
- A path is traced from the root to a leaf node which holds the prediction for that tuple

► Example

RID	age	income	student	credit-rating	Class
1	youth	high	no	fair	?

- Test on age: youth
- Test of student: no
- Reach leaf node
- **Class NO:** the customer Is Unlikely to buy a computer



A decision tree indicating whether a customer is likely to purchase a computer

4.2 Decision Tree Induction

► Why decision trees classifiers are so popular?


- The construction of a decision tree does not require any domain knowledge or parameter setting
- They can handle high dimensional data
- Intuitive representation that is easily understood by humans
- Learning and classification are simple and fast
- They have a good accuracy

► Note

- Decision trees may perform differently depending on the data set

► Applications

- medicine, astronomy
- financial analysis, manufacturing
- and many other applications



A decision tree indicating whether a customer is likely to purchase a computer

4.2.1 The Algorithm

Principle

- Basic algorithm (adopted by ID3, C4.5 and CART): a **greedy algorithm**
- Tree is constructed in a top-down recursive divide-and-conquer manner


► Iterations

- At start, all the training tuples are at the root
- Tuples are partitioned recursively based on selected attributes
- Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)

► Stopping conditions


- All samples for a given node belong to the same class
- There are no remaining attributes for further partitioning – **majority voting** is employed for classifying the leaf
- There are no samples left

Example



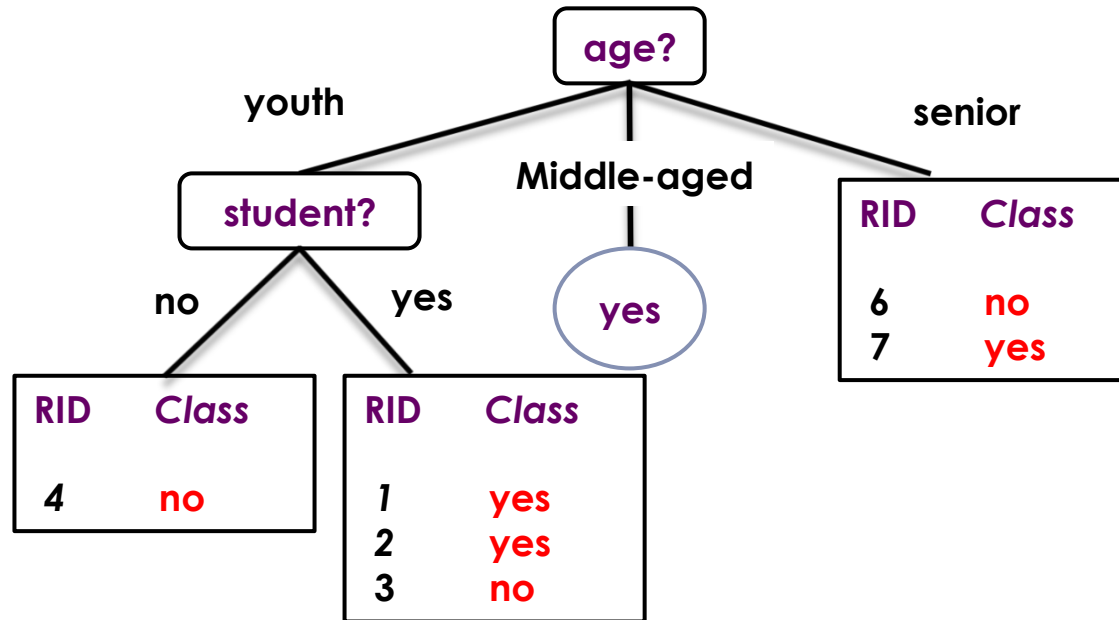
RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	no
7	senior	yes	excellent	yes

Example



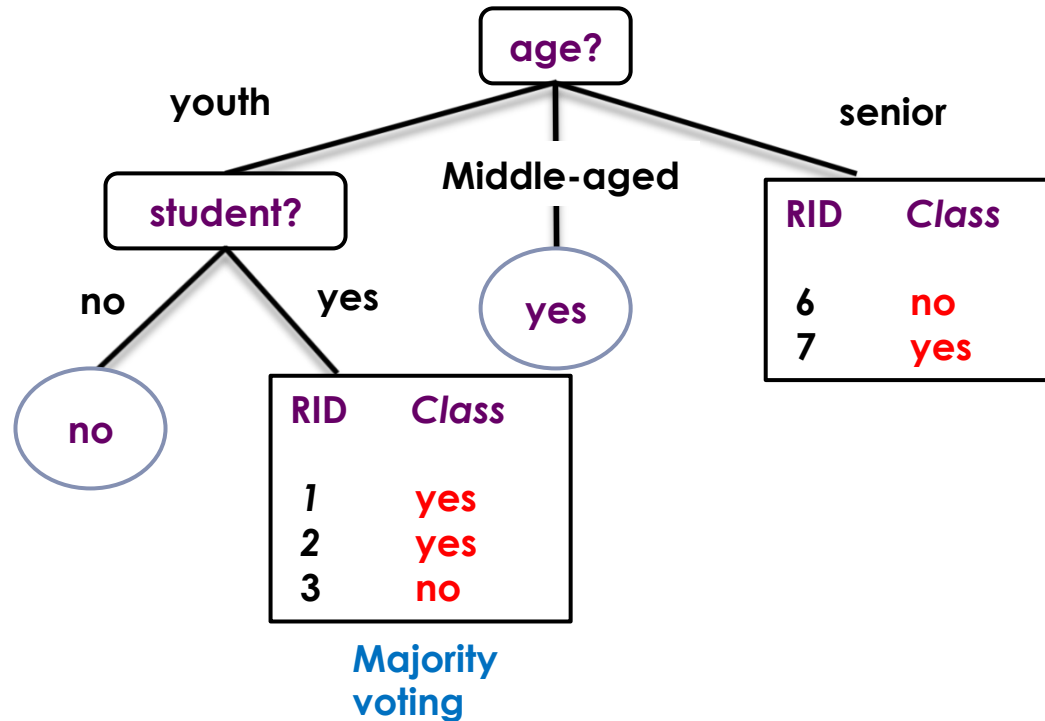
RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	no
7	senior	yes	excellent	yes

Example




RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	no
7	senior	yes	excellent	yes

Example



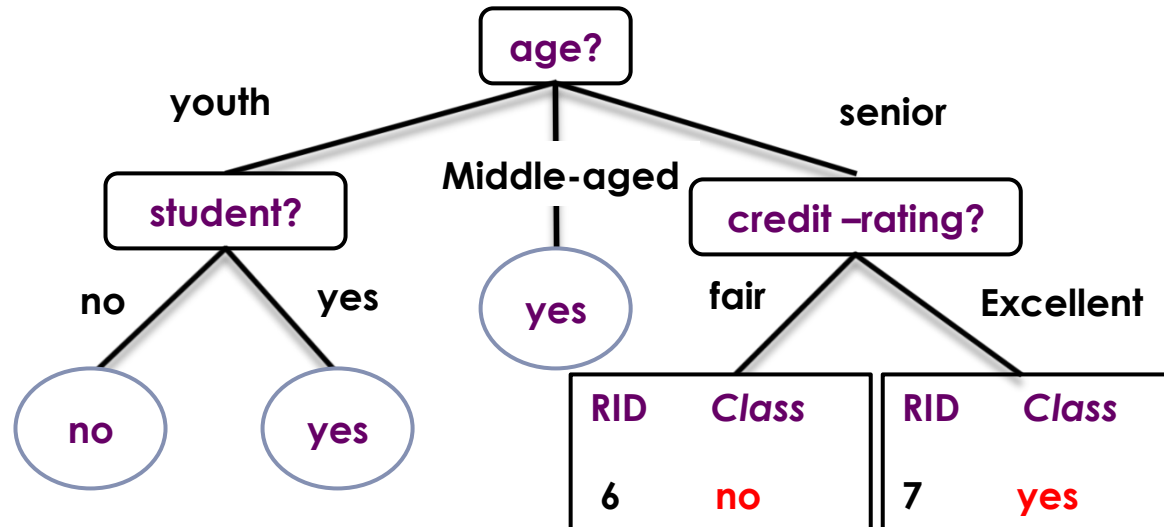
RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	no
7	senior	yes	excellent	yes

Example




RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	no
7	senior	yes	excellent	yes

Example



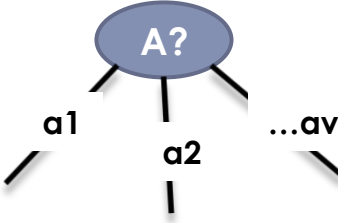

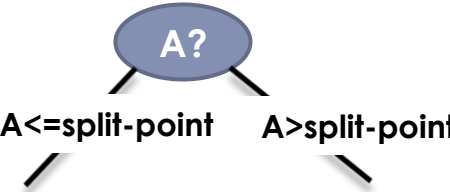
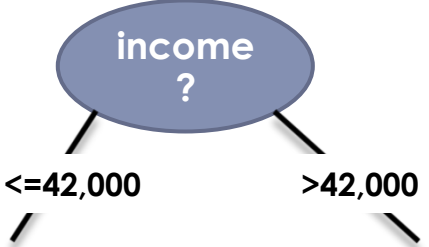
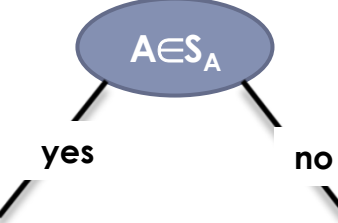

RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	no
7	senior	yes	excellent	yes

Example



RID	age	student	credit-rating	Class: buys_computer
1	youth	yes	fair	yes
2	youth	yes	fair	yes
3	youth	yes	fair	no
4	youth	no	fair	no
5	middle-aged	no	excellent	yes
6	senior	yes	fair	yes
7	senior	yes	excellent	no

Three Possible Partition Scenarios

Partitioning scenarios	Examples
<p>Discrete-valued</p> 	
<p>Continuous-valued</p> 	
<p>Discrete-valued+ binary tree</p> 	

4.2.2 Attribute Selection Measures

- ▶ An **attribute selection measure** is a heuristic for selecting the splitting criterion that “best” separates a given data partition **D**

Ideally

- Each resulting partition would be pure
- A **pure** partition is a partition containing tuples that all belong to the same class
- ▶ Attribute selection measures (splitting rules)
 - Determine how the tuples at a given node are to be split
 - Provide ranking for each attribute describing the tuples
 - The attribute with highest score is chosen
 - Determine a **split point** or a **splitting subset**
- ▶ **Methods**
 - Information gain
 - Gain ratio
 - Gini Index

Before Describing Information Gain

Entropy & Bits

You are watching a set of independent random sample of X

- ▶ X has 4 possible values:

$$P(X=A)=1/4, P(X=B)=1/4, P(X=C)=1/4, P(X=D)=1/4$$

- ▶ You get a string of symbols ACBABBCDADDC...
- ▶ To transmit the data over binary link you can encode each symbol with bits ($A=00$, $B=01$, $C=10$, $D=11$)
- ▶ You need 2 bits per symbol

Before Describing Information Gain

Fewer Bits – example 1

- ▶ Now someone tells you the probabilities are not equal

$$P(X=A)=1/2, P(X=B)=1/4, P(X=C)=1/8, P(X=D)=1/8$$

- ▶ Now, it is possible to find coding that uses only 1.75 bits on the average. How?

→ E.g., Huffman coding

Before Describing Information Gain

Fewer Bits – example 2

- ▶ Suppose there are three equally likely values

$$P(X=A)=1/3, P(X=B)=1/3, P(X=C)=1/3$$

- ▶ Naïve coding: A = 00, B = 01, C=10
- ▶ Uses 2 bits per symbol
- ▶ Can you find coding that uses 1.6 bits per symbol?
- ▶ In theory it can be done with 1.58496 bits

Before Describing Information Gain

Entropy – General Case

- ▶ Suppose X takes n values, V_1, V_2, \dots, V_n , and

$$P(X=V_1)=p_1, P(X=V_2)=p_2, \dots, P(X=V_n)=p_n$$

- ▶ What is the smallest number of bits, on average, per symbol, needed to transmit the symbols drawn from distribution of X ? It's

$$H(X) = - \sum_{i=1}^m p_i \log_2(p_i)$$

- ▶ $H(X)$ = the **entropy** of X

Before Describing Information Gain

Entropy is a measure of the average information content one is missing when one does not know the value of the random variable

► High Entropy

- X is from a **uniform** like distribution
- Flat histogram
- Values sampled from it are less predictable

► Low Entropy

- X is from a varied (**peaks and valleys**) distribution
- Histogram has many lows and highs
- Values sampled from it are more predictable

1st approach: Information Gain Approach

D: the current partition

N: represent the tuples of partition D

- ▶ Select the attribute with the highest information gain (based on the work by Shannon on information theory)
- ▶ This attribute
 - minimizes the information needed to classify the tuples in the resulting partitions
 - reflects the least randomness or “impurity” in these partitions
- ▶ **Information gain** approach minimizes the expected number of tests needed to classify a given tuple and guarantees a simple tree

Information Gain Approach

Step1: compute **Expected information** (entropy) of **D -Info(D)-**

- ▶ The expected information needed to classify a tuple in **D** is given by:

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$

- **m**: the number of classes
 - **p_i**: the probability that an arbitrary tuple in **D** belongs to class **C_i**
estimated by: **|C_{i,D}| / |D|** (proportion of tuples of each class)
 - A **log** function to the base 2 is used because the information is encoded in bits
- ▶ **Info(D)**
- The average amount of information needed to identify the class label of a tuple in D
 - It is also known as **entropy**

Info(D): Example

RID	age	income	student	credit-rating	class:buy_computer
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle-aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle-aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle-aged	medium	no	excellent	yes
13	middle-aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

m=2 (the number of classes)

N= 14 (number of tuples)

9 tuples in class yes

5 tuples in class no

The entropy (Info(D)) of the current partition D is:

$$Info(D) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940 \text{ bits}$$

Information Gain

Step2: for each attribute, compute the amount of information needed to arrive at an exact classification after partitioning using that attribute

- ▶ suppose that we were to partition the tuples in D on some attribute **A** $\{a_1, \dots, a_v\}$

- If A discrete: v outcomes

- D** split into **v** partitions $\{D_1, D_2, \dots, D_v\}$

- **Ideally** D_i partitions are pure but it is unlikely

- This amount of information we still need to arrive at an exact classification is measured by:

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

- $|D_i| / |D|$: the weight of the jth partition

- **Info(D_j)**: the entropy of partition D_j

- The smaller the expected information still required, the greater the purity of the partitions

Info_{age}(D): Example

RID	age	income	student	credit-rating	class:buy_computer
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle-aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle-aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle-aged	medium	no	excellent	yes
13	middle-aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

Using attribute age

1st partition (youth) **D1** has **2 yes** and **3 no**

$I(2,3)$ = entropy of D1(Info(D1))

2nd partition (middle-aged) **D2** has **4 yes** and **0 no**

$I(4,0)$ = entropy of D2(Info(D2))

3rd partition (senior) **D3** has **3 yes** and **2 no**

$I(3,2)$ = entropy of D3(Info(D3))

Info_{age}(D) is:

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

Information Gain Approach

- ▶ **Step1:** compute **Expected information** (entropy) of the current partition **Info(D)**
- ▶ **Step2:** compute **Info_A(D)**, the amount of information would we still need to arrive at an exact classification after partitioning using attribute A
- ▶ **Step3:** compute information gain:
- ▶ Information gain by branching on A is

$$Gain(A) = Info(D) - Info_A(D)$$

- ▶ **Information gain** is the expected reduction in the information requirements caused by knowing the value of A
 - The attribute **A** with the **highest information gain**, $(Gain(A))$, is **chosen** as the splitting attribute at node N

Info_{age}(D): Example

RID	age	income	student	credit-rating	class:buy_computer
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle-aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle-aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle-aged	medium	no	excellent	yes
13	middle-aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

$$1) \text{ Info}(D) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940 \text{ bits}$$

$$2) \text{ Info}_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$$3) \text{ Gain}(age) = \text{Info}(D) - \text{Info}_{age}(D) = 0.246$$

Similarly, Gain(Income)=0.029, Gain(student)=0.151, Gain(credit_rating)=0.48

Attribute age has the highest gain \Rightarrow It is chosen as the splitting attribute

Note on Continuous Valued Attributes

- ▶ Let attribute A be a continuous-valued attribute
- ▶ Must determine the **best split point** for A
 - Sort the value A in increasing order
 - Typically, the midpoint between each pair of adjacent values is considered as a possible *split point*
 - $(a_i + a_{i+1})/2$ is the midpoint between the values of a_i and a_{i+1}
 - The point with the *minimum expected information requirement* for A is selected as the split-point for A
- ▶ **Split**
 - **D_1** is the set of tuples in D satisfying **$A \leq \text{split-point}$** , and **$D_2$** is the set of tuples in D satisfying **$A > \text{split-point}$**